




The Brilliant Employee · 03

My routing log never recorded what ran

A triage step sized 89% of 3,717 routing decisions to the cheapest model, and I quoted that for five weeks as though it described what ran. It only ever recorded what it advised. Of the 816 jobs sized cheap, the flagship ran 796.

sgnk.ai/brilliant ↗

Sagnik Mitra

Why this exists

A language model writes a large share of my production code. Wrong work reads exactly like right work, and no prompt fixes that. So the system below exists to find out when the answers were wrong.

From one week, none of it visible from inside the tool: two fixes written up as done and never made, and a detector whose 1,112 catches were mostly one test string of mine.

THE SYSTEM, AS OF 7 AUGUST 2026, 22:55Z	
rules written	11 June, 57 days ago
ledger running	27 June, 41 days
tasks in the ledger	3,140, across 3,680 rows
controls on tool calls	6 registered, 4 that can refuse

You can rent the model. You cannot rent the part that tells you it was wrong.

Five parts, and where today sits

Five parts sit around the model. Each post takes one of them.

sizing check	sizes a job before a model is picked
---------------------	--------------------------------------

model picker	recommends which model gets the job
---------------------	-------------------------------------

safety stops	refuse the irreversible, before it runs
---------------------	---

grader	grades finished work pass or fail
---------------	-----------------------------------

task log	at least one append-only line per task
-----------------	--

Listed in the order they act on a task, not the order they were built.

Today is about the sizing check, marked above.

ASIDE. Only the safety stops can refuse. The other four observe: one advises, one recommends, one grades after the fact, one records.

This page is the map. Nothing here needs another post to make sense.

Every task is sized before a model sees it

Nothing goes straight to a model. Every request meets the sizing check first: is this trivial, does it need the expensive model, or is the scope genuinely unenumerable.

Plain words: SIZING CHECK. a pre-flight triage step that runs before any model call. It sizes the task and recommends a tier, plus whether the job needs verification or approval. It records all of that and enforces none of it.

- a request arrives
typed by me, or raised by a job already running
- the sizing check sizes it and writes one line
which tier, one shot or a workflow, guarded or not
- something else picks the model and runs it
a separate component, which that line cannot stop

Its last line is exit zero.

The standing rule behind it is blunt: the best model call is the one you never make. Trivial work skips the whole orchestration stack.

Triage before assignment. Always.

Most work is small, so most work goes cheap

Cheap and single-shot

the floor model does the job alone. 89% of recommendations land here

Strong model

reserved for architecture, multi-file surgery, and deep investigation. 11% of recommendations.

Full workflow

multiple coordinated agents, only for scope you cannot enumerate up front. 5%.

The split is not thrift. It is the shape the work actually has. Most tasks are small, and pretending otherwise burns tokens on ceremony while teaching you nothing about which tasks were ever hard.

ASIDE. Note the word recommendations. It is doing more work in that sentence than I realised when I wrote it.

Cheap by default. Expensive by exception.

89% cheap is the design, not the corner cut

THE SIZING CHECK'S ADVICE, PINNED TO 4 AUGUST 22:48Z	
recommended floor (cheap)	3,311 = 89%
recommended strong	406 = 11%
recommended single shot	3,547 = 95%
recommended workflow	170 = 5%
floor AND guarded	861 = 23%
every row is an intention. Not one of them is an outcome.	

Most people read 89% cheap as corner-cutting. It is the design. The highlighted row is the one that says whether the design is working.

ASIDE. A distribution over advice is not a distribution over spend. For five weeks I read this table as though it described what my system did. It describes what my system advised.

A recommendation is not an outcome.

The sizing check writes a line and exits zero

602

2,064

of the first 2,064 decisions, flagged as needing verification or explicit approval, counted across every tier. Those rows end on 28 July

Guarded means the sizing check appended a sentence to the worker's instructions: run the verify ladder before calling this done, or name this operation and get approval before running it. That is the whole mechanism. It writes the line and exits zero.

The word guard oversells it. Nothing is blocked at this step; enforcement is a different component. What this contributes is that the instruction rides along automatically instead of depending on me remembering.

Name what your guard actually does, not what you hoped it would.

The old target measured cost, not safety

OLD TARGET, RETIRED 28 JULY

metric

% blocked from the heavy path

target

30 to 50%

problem

measures cost, not safety

CURRENT TARGET

default

floor tier, single shot

strong tier

architecture, multi-file, deep digs

workflow

unenumerable scope only

The old target died because the sizing check stopped emitting the thing it measured.

23.2%

the guarded fraction at the cutoff, 861 of 3,717, against 23.3% over the first 2,064 decisions. Flat.

Watch the guardrail metric, not the cost metric.

There is no field for which model ran

I went looking for the model field, to confirm the 89% held in practice.

```
$ jq -r --arg c "$CUT" 'select(
  .gate_tier=="floor" and
  .timestamp<=$c).model' \
  traces/*.jsonl | sort | uniq -c
698 claude-opus-5
98 claude-opus-4-8
10 unknown
8 claude-fable-5
2 claude-sonnet-5
```

2 of 816

ledger rows carrying a floor-tier gate verdict ran on the cheap tier. The flagship ran 796 of them, and 10 name no model at all. Only 915 of 3,357 rows carry a gate verdict.

A log that records intent can never tell you it is being ignored.

The best model for everything measures nothing

REJECTED. A router that sends every task to the flagship model. It feels safe and measures nothing: you pay for reputation, collect no verification record, and never learn which tasks were trivial all along. I rejected this design in June, built the sizing check to avoid it, and then ran it anyway, unnoticed, across every one of those 3,717 decisions.

That is the part worth taking. The design was right. The measurement was pointed at the wrong file.

FLAGSHIP FOR EVERYTHING

you learn

nothing

you pay

every task

CHEAP PLUS A CHECKER

you learn

where it

breaks

you pay

on escalation

Reliability lives in the workplace, not the hire.

Log the outcome, or you are logging your own opinion.



Cheap by default. Expensive by exception. Guarded on purpose.

What 2,064 live gate decisions actually look like, and the target I retired.

sgnk.ai/brilliant ↗

Sagnik Mitra

I build AI systems and the machinery that keeps them honest: gates, ledgers, judges, and rails. Product design for years, engineering the whole time. This series is what the building actually taught me.

sgnk.ai · in/sagnikmitra · github.com/sagnikmitra