




The Brilliant Employee · 01

My AI agents committed to git anyway

A language model writes most of my production code, and it fails in the same voice it succeeds in. So the checks moved into programs on the tool boundary: six registered, four that can refuse.

sgnk.ai/brilliant ↗

Sagnik Mitra

Why this exists

A language model writes a large share of my production code. Wrong work reads exactly like right work, and no prompt fixes that. So the system below exists to find out when the answers were wrong.

From one week, none of it visible from inside the tool: a sizing step ignored on 796 of 816 tasks, and two fixes written up as done that were never made.

THE SYSTEM, AS OF 7 AUGUST 2026, 14:51Z	
rules written	11 June, 57 days ago
ledger running	27 June, 41 days
tasks logged	3,103, across 3,620 rows
controls on tool calls	6 registered, 4 that can refuse

You can rent the model. You cannot rent the part that tells you it was wrong.

Five parts, and where today sits

Five parts sit around the model. Each post takes one of them.

sizing check	sizes a job before a model is picked
model picker	recommends which model gets the job
safety stops	refuse the irreversible, before it runs
grader	grades finished work pass or fail
task log	at least one append-only line per task

Listed in the order they act on a task, not the order they were built.

Today cuts across more than one, so no row is marked.

ASIDE. Only the safety stops can refuse. The other four observe: one advises, one recommends, one grades after the fact, one records.

This page is the map. Nothing here needs another post to make sense.

No interview predicts what this model does next

The worker is software: a language model that writes production code in my studio. One afternoon it came close to deleting a client's business, then reported it in the same calm tone it uses for everything else.

The part worth sitting with is not the mistake. It is the tone. The same worker who moves a client's whole database before lunch will invent a setting that does not exist after it, and both answers arrive equally sure of themselves.

An interview samples a few answers and trusts them to stand for the whole. Nothing here holds still long enough for that to work.

THE WORKER, IN PLAIN TERMS	
what it is	a language model in a loop
what it can do	read files, run commands, call APIs
what it remembers	nothing, once the session ends
how it fails	fluently, and at the same speed

Sampling behaviour once does not bound behaviour later.

A written rule cannot refuse

What I got wrong first was assuming that writing a rule down was the same as enforcing it.

Everything you tell a language model is text. A capitalised DO NOT is not a different kind of object: it lands in the same space as the question, and sometimes it loses the argument.

Plain words: A CONTROL. An ordinary program that runs before the action does and can refuse it. There is nothing in it to persuade, because it is not made of language. A git pre-commit hook is the same idea.

```
~/ .sgnk/bin/sgnk-bash-guard.sh      header
# By the system's own LR#48 (a textual prohibition
# is advisory, never the gate)
```

The first clause of the guard's header comment. The lesson behind it was expensive.

ASIDE. None of this shape is new and I did not invent it. seccomp does the same job for system calls, at the kernel boundary.

**A rule the model reads is advice. A
program on the boundary is a control.**

Six controls registered. Four can refuse.

A control is a script registered against a tool, and it runs first. Three of mine refuse by exiting with a specific code. One refuses by returning a permission decision instead. Two only watch.

```
~/.claude/settings.json    hooks.PreToolUse
{
  "matcher": "Bash",
  "hooks": [{
    "type": "command",
    "command": "bash ~/.sgnk/bin/sgnk-bash-guard.sh",
    "timeout": 5
  }]
}
```

matcher picks the tool. command is the program that gets to say no. The home directory is shortened to a tilde here; the file carries the absolute path.

Six sit on tool calls: shell, file edits, skills, agent spawning, and two on anything reaching outside. Only three are armed. Two can never stop anything by contract, and a third waits behind an unset flag.

Registered is not armed.

Two of seven steps can refuse

- I type a sentence.
on submit
- The job is sized before a model is picked.
the sizing check
- A model is recommended, cheapest first.
the model picker
- It asks to run a command.
the safety stops | catastrophic only,
refuse
- It asks to read something off the web.
the safety stops | untrusted source,
refuse
- The work is graded pass or fail, by a grader I do not fully trust.
the grader
- One line is written down, and never edited.
the task log

28 hooks across 8 moments in a session. The sizing check is fail-open on purpose: a triage nudge that can wedge a turn is worse than no triage. The grader is graded too, and the one calibration of the sizing check scored 0.1 against a bar of 0.7.

Three of those five parts contain no language model at all.

Two of the seven steps can refuse, and both only for a short list.

541 warnings, and none blocked

Here is the part I would skip if I were selling you something. My shell guard has two tiers, and the wide one does not block. It logs the hit, names the rule I am about to walk into, and lets me through.

```
~/sgnk/bin/sgnk-bash-guard.sh      the always-block list  
git push .*--force  
DROP (TABLE|DATABASE|SCHEMA) | TRUNCATE | dropdb  
vercel (rm|remove|rollback)  
aws .* (delete|terminate|remove|rm)
```

Four of ten, abbreviated. This list exits 2 whatever mode the guard is in, though three doors sit above it: a bypass variable, a one-shot human approval token, and a recursion guard for nested runs.

Everything wider only warns, and that is the lesson. A gate that blocks everything gets bypassed within a week, and a bypassed gate protects nothing.

ASIDE. Counted between 11 July and 4 August. The guard matches the text of a command, not the intent behind it, so grepping its own deny-list counts as a hit. It flagged one of mine while I wrote this slide.

**A gate that blocks everything gets
switched off. Then you have no gate.**

A thumbs up is not an approval

The heaviest of the eight standing rules governs anything irreversible. It does not say be careful. It specifies a ritual, and the ritual is what makes it checkable.

- 1 Name the operation in plain language.
- 2 State the blast radius. How many rows, what breaks, when the last backup ran.
- 3 Say what was checked to confirm it is safe.
- 4 Ask, and then stop.
- 5 Wait for an explicit yes.
- 6 Do it, then verify the state afterwards and report it.

Every operation is its own approval. Yesterday's yes does not carry.

ASIDE. Ok, continue, a thumbs up, and silence are each written into the rule as not-approval. Every one of them had to be listed, because every one of them had been tried.

Only a fresh yes counts.

Every finished task leaves one line

A memory the model can edit is not evidence. So the record lives outside it: one line appended per finished task, in a file nothing rewrites. This is a real line from the session that produced this deck.

```
~/sgnk/traces/2026-08-04.jsonl
"gate_tier":      "floor",
"decided_model":  "sonnet",
"decision_mode":  "advisory",
"decision_real_n": 0,
"model":          "claude-opus-5",
```

Five of thirty fields, reindented to fit. Values unchanged.

Read in order, it indicts itself. The cheap model was recommended. The expensive one ran anyway, because the recommendation is advisory and nothing enforces it. And the count of live observations behind that recommendation was zero, which is what a system looks like before it has learned anything.

One line per finished task since 27 June.
47 of them carry a verdict I gave myself.

68 corrections, none removed

Seven rules govern the work. The eighth governs the rulebook. Every correction, mine or one the system catches in its own work, becomes a numbered line in a second block that may only ever be added to.

```
~/.claude/CLAUDE.md      Learned Rules
```

```
66. Run every harness assertion under BOTH bash  
    and zsh before trusting its verdict.
```

```
67. A tool that writes must verify the write  
    LANDED before reporting success.
```

```
68. A test whose subject is a rare fault proves  
    nothing until it reproduces the fault.
```

The last three, abbreviated to their first sentence. 67 was written the day a tool printed REGISTERED twice while storing nothing.

The eight rules are still eight. The block above them stands at 68, with no gaps. It is the one thing here that accumulates.

ASIDE. The system reads it at the start of every session and forgets everything else.

The original eight are unchanged since June. The 68 above them each came from one incident.



I never interviewed the model, so every task gets checked.

The full argument, the numbers behind it, and the commands that derive them.

sgnk.ai/brilliant ↗

Sagnik Mitra

I build AI systems and the machinery that keeps them honest: gates, ledgers, judges, and rails. Product design for years, engineering the whole time. This series is what the building actually taught me.

sgnk.ai · in/sagnikmitra · github.com/sagnikmitra