




The Brilliant Employee · 08

The test passed. I rejected the work anyway.

Every task my terminal AI finishes writes one line into a log file it cannot edit. Here is one real line from 3 August, taken apart field by field: what the model picker recommended, what actually ran, and whether I kept the work.

sgnk.ai/brilliant ↗

Sagnik Mitra

Why this exists

sgnk is a one-person studio. A language model writes a large share of the production code. You do not need your own system to get answers out of a model. You need one to find out when the answers were wrong, because wrong work reads exactly like right work and no prompt fixes that.

Three things from that week, none visible from inside the tool: a sizing step ignored on 796 of 816 tasks, two fixes I had written up as done that were never made, and a detector whose 1,112 catches were mostly one test string of my own.

THE SYSTEM, AS OF 6 AUGUST 2026, 16:43Z	
rules written	11 June, 56 days ago
ledger running	27 June, 40 days
tasks logged	3,063, across 3,562 rows
controls on tool calls	6 registered, 3 of them armed

You can rent the model. You cannot rent the part that tells you it was wrong.

Five parts, and where today sits

Five parts sit around the model. What changes from one piece of writing to the next is which of them is under the microscope.

the Gate	sizes a job before a model is picked
the Gambler	recommends which model gets the job
the Rails	refuse the irreversible, before it runs
the Judge	grades finished work pass or fail
the Ledger	at least one append-only line per task

Listed in the order they act on a task, not the order they were built.

Today is about the Ledger, marked above. The other four are the machinery around it, and this post does not assume you know any of them.

ASIDE. Only the rails can refuse. The gate advises, the router recommends, the judge grades after the fact, and the ledger only records. Four of the five observe; one intervenes.

This page is the map. Nothing here needs another post to make sense.

The worker cannot edit his own record

The Ledger is the part of the workplace that remembers. Written rule #31: every task appends ONE line to a dated file, with a correlation id and an accepted flag. Every task, not some.

```
if [ -s "$TRACE_FILE" ]; then
    printf '\n%s' "$_ROW" >> "$TRACE_FILE"
else
    printf '%s' "$_ROW" >> "$TRACE_FILE"
fi
```

The employee, my terminal AI, writes that line through a hook he cannot touch and cannot revise afterwards. A worker who grades his own homework will also tidy his own transcript.

Plain words: THE LEDGER. ~/.sgnk/traces/YYYY-MM-DD.jsonl. One file per day, at least one JSON line per task, written by the harness, never by the worker.

If the record can be edited by the thing being measured, it is not a record.

A whiteboard arrow is not a working loop

Written rule #32: the words self-improving are banned in this workspace unless a grading loop is RUNNING and feeding the model picker.

- Every finished task appends one line
the harness writes it, not the worker
- A grading skill scores that work
machine check, pass or fail
- Those grades feed the model picker | never built
the arrow on the whiteboard

Rule #32 unlocks the phrase only when that last step provably runs.

The failure mode is seductive. You build the ledger, you build the grader, you draw the arrow between them on a whiteboard, and the sentence 'the system learns from its mistakes' starts to feel true. The arrow is the part that never gets built.

ASIDE. The ban applies to me as much as to him. I wrote it after catching myself about to claim the loop was closed when it was not.

A learning claim is only as real as the pipeline that consumes the data.

What one finished task looks like

Pulled from the live file for 2026-08-03, redacted only where it names the session. One task: a generate-shaped job, floor tier, single shot, in my own Macaulay2 web project at 00:12 UTC.

```
~/sgnk/traces/2026-08-03.jsonl (one row, redacted)
```

```
{"timestamp":"2026-08-03T00:12:32Z","correlation_id":"SESSION:1...  
"skill":"unknown","gate_tier":"floor","gate_orchestrate":"singl...  
"decided_model":"sonnet","decision_real_n":0,"decision_mode":"a...  
"task_shape":"generate","model":"claude-opus-5","tier":"opus",  
"input_tokens":838575,"fresh_input_tokens":2,"cache_read_tokens...  
"output_tokens":9896,"latency_ms":191000,"cwd":"~/GitHub/m2web-...  
"assertion_pass":true,"failure_mode":null,"learning_mode":"on_p...  
"accepted":false,"assertion_n":1,"assertion_source":"machine-ga...
```

ASIDE. Trimmed to fit the slide. The full row carries 33 keys.

Everything above is one task's complete paper trail, written the moment it finished.

The cheap model was picked. The big one ran.

ROUTING FIELDS	
correlation_id	SESSION:16 . the join key every downstream loop keys on
gate_tier	floor . the Gate, the part that sizes the job, sent this to the cheap tier
decided_model	sonnet . the Gambler, the part that recommends the model, recommended cheap
decision_mode	advisory . the pick was logged, not enforced. The big model ran anyway
decision_real_n	0 . real observations behind that recommendation: none
The row keeps the gap between recommendation and reality.	

Both parts spoke, and the line preserves what each said next to what actually happened. Decided sonnet, ran opus. That gap is signal you can only collect by writing both halves down.

Log the recommendation AND the outcome.
The delta between them is the signal.

Two verdicts, one line, joined forever

COST AND VERDICTS	
input_tokens	838,575 . of which 838,573 came from cache. Fresh tokens: 2
output_tokens	9,896
latency_ms	191,000 . three minutes eleven seconds of work
assertion_pass	true . the Judge, the machine check that grades, passed it 1 of 1
accepted	false . I rejected the work anyway
Machine says yes, human says no. That gap is how a miscalibrated grader gets caught.	

Rare, and the reason the line exists. Of 3,357 ledger rows, 38 carry both a machine verdict and a human one. 18 agree, 20 disagree, and 6 disagree in this direction: machine pass, human reject.

A pass from the Judge plus a no from me is worth more than ten clean agreements.

Each field has a reader. Not all read yet.

- 1 accepted feeds the preference log: every accept and reject stored as a pair, per skill and model
- 2 assertion_pass feeds reward mining, joined to accepted through the correlation_id
- 3 decided_model versus what ran feeds the scorecard of the part that recommends the model, once real observations exist
- 4 drift watch compares this week's rows against snapshot baselines
and the rulebook only unlocks the phrase self-improving when these consumers provably run

That is the design. The honest status is printed on the line itself: decision_mode says advisory, decision_real_n says 0. The wiring exists, the current is intermittent.

ASIDE. My routing brain writes its homework down and does not yet read it. Which is precisely why the ban in rule #32 has not been lifted.

The ledger is complete. The reading of it is the unfinished part, and the line admits that.

One line per task is the cheapest thing here

39

daily files, one per day from 2026-06-27 to 2026-08-04, no missing dates.

Because every loop I want needs this exact substrate. You cannot calibrate a grader without machine grades joined to human verdicts. You cannot route by track record without a track record. You cannot detect drift without a baseline made of dated rows.

One line per task, and every expensive thing built since sits on top of it. Every claim I am allowed to make in public is bounded by what this file can prove.

REJECTED. A dashboard-first analytics stack. The line came first; charts wait for consumers that actually exist.

Measurement before improvement. The order is not negotiable.



A machine check passed this task. I rejected it anyway. Both verdicts sit on one line, in a file the employee never writes himself. The full dissection, and the commands to check it yourself.

sgnk.ai/brilliant ↗

Sagnik Mitra

I build AI systems and the machinery that keeps them honest: gates, ledgers, judges, and rails. Product design for years, engineering the whole time. This series is what the building actually taught me.

sgnk.ai · in/sagnikmitra · github.com/sagnikmitra